

# Zang and Wang 2019: Neural Dynamics on Complex Networks

Minqi Pan

April 4, 2020

# Neural Dynamics on Complex Networks

- AAI 2020, Best Paper of “The 1st International Workshop on Deep Learning on Graphs: Methodologies and Applications”, Feb 8th, 2020
- Chengxi Zang and Fei Wang
- Weill Cornell Medicine

# Outline

- 1 General Framework
  - Neural Dynamics on Complex Networks (NDCN)
- 2 Learning Continuous-Time Network Dynamics
  - Model Instance
  - Experiments
- 3 Learning Regularly-Sampled Dynamics
  - Baselines, Experimental Setup and Results
- 4 Learning Semantic Labels at Terminal Time
  - Model Instance
  - Experiments

# Outline

- 1 General Framework
  - Neural Dynamics on Complex Networks (NDCN)
- 2 Learning Continuous-Time Network Dynamics
  - Model Instance
  - Experiments
- 3 Learning Regularly-Sampled Dynamics
  - Baselines, Experimental Setup and Results
- 4 Learning Semantic Labels at Terminal Time
  - Model Instance
  - Experiments

# The Differential Equation System

$$\frac{dX(t)}{dt} = f(X(t), G, W(t), t)$$

- $X(t) \in \mathbb{R}^{n \times d}$ : the state (node feature values) of a dynamic system consisting of  $n$  linked nodes at time  $t \in [0, \infty)$ , and each node is characterized by  $d$  dimensional features
- $f: \mathbb{R}^{n \times d} \rightarrow \mathbb{R}^{n \times d}$ : a function governing the dynamics of the system, which could be either linear or nonlinear
- $G = (\mathcal{V}, \mathcal{E})$ : the network structure capturing how the nodes are linked to each other
- $W(t)$ : the parameters which control how the system evolves over time
- $X(0) = X_0$ : the initial state of this system at time  $t = 0$

# Semantic Labels

- $Y(X, \Theta, t) \in \{0, 1\}^{n \times k}$ : the semantic labels of the nodes at time  $t$
- $\Theta$ : the parameters of this classification function

# Problem #1: Network Dynamics Learning

- Given a graph  $G$  and the observations of the states of system:

$$\{X(\hat{t}_1), X(\hat{t}_2), \dots, X(\hat{t}_T) : 0 \leq t_1 \leq \dots \leq t_T\}$$

- $t_1$  to  $t_T$  are arbitrary physical time stamps, possibly irregularly sampled with different observational time intervals
- How to learn the continuous-time dynamics  $\frac{dX(t)}{dt}$  on complex networks from empirical data? Can we learn differential equation systems  $\frac{dX(t)}{dt} = f(X(t), G, W(t), t)$  to generate or predict continuous-time dynamics  $X(t)$  at arbitrary physical time  $t$ ?
  - “extrapolation prediction”: when  $t > t_T$
  - “interpolation prediction”: when  $t < t_T$  and  $t \neq \{t_1, \dots, t_T\}$

## Problem #2: Structured Sequence Learning

- A special case of the problem of Network Dynamics Learning
- $t_1, t_2, \dots, t_T$  are sampled regularly with equal time intervals
- Emphasizing on sequential order instead of arbitrary physical time
- The goal is to extrapolate next  $m$  steps:

$$X[t_T + 1], \dots, X[t_T + m]$$



## Problem #3: One-snapshot Learning

- A special case of the problem of Network Dynamics Learning
- How to learn the semantic labels of  $Y(X(t_T))$  at the moment  $t = t_T$  for each node?
- Emphasizing on a specific moment
- Without loss of generality, we focus on the moment at the terminal time  $t_T$
- The function  $Y$  can be a mapping from the nodes' states (e.g. humidity) to their labels (e.g. taking umbrella or not)

# Network Dynamics #1: Heat Diffusion

- Let  $\overrightarrow{x_i(t)} \in \mathbb{R}^{d \times 1}$  be  $d$  dimensional features of node  $i$  at time  $t$
- Thus

$$X(t) = \begin{bmatrix} \vdots \\ \overrightarrow{x_i(t)} \\ \vdots \end{bmatrix}$$

- The heat diffusion dynamics governed by Newton's law of cooling

$$\frac{d\overrightarrow{x_i(t)}}{dt} = -k_{i,j} \sum_{j=1}^n A_{i,j} (\overrightarrow{x_i} - \overrightarrow{x_j})$$

which states that the rate of heat change of node  $i$  is proportional to the difference of the temperature between node  $i$  and its neighbors with heat capacity matrix  $A$

## Network Dynamics #2: Mutualistic Interaction

- The mutualistic differential equation systems capture the abundance  $\vec{x}_i(t)$  of species  $i$  in ecology:

$$\frac{d\vec{x}_i(t)}{dt} = b_i + \vec{x}_i \left( 1 - \frac{\vec{x}_i}{k_i} \right) \left( \frac{\vec{x}_i}{c_i} - 1 \right) + \sum_{j=1}^n A_{i,j} \frac{\vec{x}_i \vec{x}_j}{d_i + e_i \vec{x}_i + h_j \vec{x}_j}$$

- with incoming migration term  $b_i$
  - with logistic growth with population capacity  $k_i$
  - with Allee effect with cold-start threshold  $c_i$
  - with mutualistic interaction term with interaction network  $A$
- For brevity, the operations between vectors are element-wise

# Network Dynamics #3: Gene Regulatory

- Governed by Michaelis-Menten equation

$$\frac{d\vec{x}_i(t)}{dt} = -b_i \vec{x}_i(t)^f + \sum_{j=1}^n A_{i,j} \frac{\vec{x}_j^h}{\vec{x}_j^h + 1}$$

- the 1st term models degradation when  $f = 1$  or dimerization when  $f = 2$
- the 2nd term captures genetic activation tuned by the Hill coefficient  $h$

# Complex Networks

- ① “Grid” where each node is connected with 8 neighbors
- ② “Random” generated by Erdős and Rényi model
- ③ “Power-law” generated by Albert-Barabási model
- ④ “Small-world” generated by Watts-Strogatz model
- ⑤ “Community” generated by random partition model

# Visualization

- To visualize dynamics on complex networks over time is not trivial
- We first generate a network with  $n$  nodes by aforementioned network models
- The nodes are re-ordered according to the community detection method by Newman
- Each node has a unique label from 1 to  $n$
- We layout these nodes on a 2-dimensional  $\sqrt{n} \times \sqrt{n}$  grid and each grid point  $(r, c) \in \mathbb{N}^2$  represents the  $i^{\text{th}}$  node where  $i = r\sqrt{n} + c + 1$
- Thus, nodes' states  $X(t) \in \mathbb{R}^{n \times d}$  at time  $t$  when  $d = 1$  can be visualized as a scalar field function  $X : \mathbb{N}^2 \rightarrow \mathbb{R}$  over the grid

# General Framework

$$\arg \min_{W(t), \Theta(T)} \mathcal{L} = \int_0^T \mathcal{R}(X(t), G, W, t) dt + \mathcal{S}(Y(X(T), \Theta))$$

subject to  $\frac{dX(t)}{dt} = f(X(t), G, W, t), X(0)$

- $\mathcal{R}(X(t), G, W, t)$ : the running loss of the dynamics on graph at time  $t$
- $\mathcal{S}(Y(X(T), \Theta))$ : the terminal semantic loss at time  $T$
- By integrating  $\frac{dX(t)}{dt} = f(X(t), G, W, t)$  over time  $t$  from initial state  $X_0$ , a.k.a. solving the initial value problem for this differential equation system, we can get the continuous-time dynamics  $X(t) = x(0) + \int_0^T f(X(\tau), G, W, \tau) d\tau$  at arbitrary time moment  $t > 0$

# As an Optimal Control Problem

- By solving the above optimization problem
  - Obtain the best control parameters  $W(t)$  for differential equation system  $\frac{dX}{dt} = f(X, G, W, t)$
  - Obtain the best classification parameters  $\Theta$  for semantic function  $Y(X(t), \Theta)$
- Differences from the traditional Optimal Control framework:  
We model the differential equation systems

$$\frac{dX}{dt} = f(X, G, W, t)$$

by graph neural networks



# In a Dynamical System View

- By integration  $\frac{dX}{dt} = f(X, G, W, t)$  over continuous time, namely

$$X(t) = X(0) + \int_0^t f(X(\tau), G, W, \tau) d\tau$$

we get our differential deep learning models

- Our differential deep learning models can be a time-varying coefficient dynamical system where  $W(t)$  changes over time
- Or a constant coefficient dynamical system when  $W$  is constant over time for parameter sharing

## Further Encoding (1)

$$\arg \min_{W(t), \Theta(T)} \mathcal{L} = \int_0^T \mathcal{R}(X(t), G, W, t) dt + \mathcal{S}(Y(X(T), \Theta))$$

subject to  $X_h(t) = f_{\text{encode}}(X(t))$

$$\frac{dX_h(t)}{dt} = f(X_h(t), G, W, t), X_h(0)$$

$$X(t) = f_{\text{decode}}(X_h(t))$$

- To further increase the express ability of our model, we can encode the network signal  $X(t)$  from the original space to  $X_h(t)$  in hidden space (usually with a different number of dimensions), and learn the dynamics in such a space

## Further Encoding (2)

$$\arg \min_{W(t), \Theta(T)} \mathcal{L} = \int_0^T \mathcal{R}(X(t), G, W, t) dt + \mathcal{S}(Y(X(T), \Theta))$$

$$\text{subject to } X_h(t) = f_{\text{encode}}(X(t))$$

$$\frac{dX_h(t)}{dt} = f(X_h(t), G, W, t), X_h(0)$$

$$X(t) = f_{\text{decode}}(X_h(t))$$

- The 1st constraint transforms  $X(t)$  into hidden space  $X_h(t)$
- The 2nd constraint is the governing dynamics in the hidden space
- The 3rd constraint decodes the hidden signal back to the original space

## Further Encoding (3)

$$\arg \min_{W(t), \Theta(T)} \mathcal{L} = \int_0^T \mathcal{R}(X(t), G, W, t) dt + \mathcal{S}(Y(X(T), \Theta))$$

$$\text{subject to } X_h(t) = f_{\text{encode}}(X(t))$$

$$\frac{dX_h(t)}{dt} = f(X_h(t), G, W, t), X_h(0)$$

$$X(t) = f_{\text{decode}}(X_h(t))$$

- The design of  $f_{\text{encode}}, f, f_{\text{decode}}$  are flexible to be any neural structure, e.g. Softmax as the decoder for classification
- We denote this model as “NDCN”

# Discrete Layers vs. Continuous Layers

- The deep learning methods with  $L$  hidden neural layers  $f_*$  are

$$X[L] = f_L \circ \cdots \circ f_2 \circ f_1(X[0]),$$

which are iterated maps with an integer number of discrete layers and thus cannot learn continuous-time dynamics  $X(t)$  at arbitrary time

- In contrast, our model

$$X(t) = X(0) + \int_0^t f(X(\tau), G, W, \tau) d\tau$$

can have continuous layers with a real number  $t$  depth corresponding to continuous-time dynamics

# Solving the Initial Value Problem

- Integrate the differential equation systems over time by numerical methods
- The numerical methods can approximate continuous-time dynamics

$$X(t) = X(0) + \int_0^t f(X(\tau), G, W, \tau) d\tau$$

at arbitrary time  $t$  accurately with guaranteed error

- In order to learn the learnable parameters  $W$ , we back-propagate the gradients of the loss function w.r.t. the control parameters  $\frac{\partial \mathcal{L}}{\partial W}$  over the numerical integration process backwards in an end-to-end manner, and solve the optimization problem by stochastic gradient descent methods

# Outline

- 1 General Framework
  - Neural Dynamics on Complex Networks (NDCN)
- 2 Learning Continuous-Time Network Dynamics
  - Model Instance
  - Experiments
- 3 Learning Regularly-Sampled Dynamics
  - Baselines, Experimental Setup and Results
- 4 Learning Semantic Labels at Terminal Time
  - Model Instance
  - Experiments

# The Continous-time Setting

- The observational times  $t_1$  to  $t_T$  of the observed states of system

$$\{X(\hat{t}_1), X(\hat{t}_2), \dots, X(\hat{t}_T) : 0 \leq t_1 \leq \dots \leq t_T\}$$

are arbitrary physical time stamps which are irregularly sampled with different observational time intervals

- Extrapolation prediction is to predict

$$X(t)$$

at arbitrary physical time moment  $t$  when  $t > t_T$

- Interpolation prediction is to predict

$$X(t)$$

when  $t < t_T$  and  $t \neq \{t_1, \dots, t_T\}$



# Model Instance (1)

$$\arg \min_{W_*, b_*} \mathcal{L} = \int_0^T |X(t) - \hat{X}(t)| dt$$

$$\text{subject to } X_h(t) = \tanh(X(t)W_e + b_e)W_0 + b_0$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t)W + b), X_h(0)$$

$$X(t) = X_h(t)W_d + b_d$$

- Loss: emphasizing on running loss only; use  $\ell_1$ -norm loss as the running loss  $\mathcal{R}$
- $|\cdot|$ :  $\ell_1$ -norm loss (element-wise absolute value difference) between  $X(t)$  and  $\hat{X}(t)$  at time  $t \in [0, T]$

## Model Instance (2)

$$\arg \min_{W_*, b_*} \mathcal{L} = \int_0^T |X(t) - \hat{X}(t)| dt$$

$$\text{subject to } X_h(t) = \tanh(X(t)W_e + b_e)W_0 + b_0$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t)W + b), X_h(0)$$

$$X(t) = X_h(t)W_d + b_d$$

- The encoding function: two fully connected neural layers with a nonlinear hidden layer as the encoding function
- the linear decoding function: for regression tasks in the original signal space

## Model Instance (3)

$$\arg \min_{W_*, b_*} \mathcal{L} = \int_0^T |X(t) - \hat{X}(t)| dt$$

$$\text{subject to } X_h(t) = \tanh(X(t)W_e + b_e)W_0 + b_0$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t)W + b), X_h(0)$$

$$X(t) = X_h(t)W_d + b_d$$

- $\hat{X}(t) \in \mathbb{R}^{n \times d}$ : the supervised dynamic information available at time stamp  $t$ 
  - in the semi-supervised case the missing information can be padded by 0

## Model Instance (4)

$$\arg \min_{W_*, b_*} \mathcal{L} = \int_0^T |X(t) - \hat{X}(t)| dt$$

$$\text{subject to } X_h(t) = \tanh(X(t)W_e + b_e)W_0 + b_0$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t)W + b), X_h(0)$$

$$X(t) = X_h(t)W_d + b_d$$

- $\Phi = D^{-\frac{1}{2}}(D - A)D^{-\frac{1}{2}} \in \mathbb{R}^{n \times n}$ : graph diffusion operator to model the instantaneous network dynamics in the hidden space, which is the normalized graph Laplacian
  - $A \in \mathbb{R}^{n \times n}$ : the adjacency matrix of the network
  - $D \in \mathbb{R}^{n \times n}$ : the corresponding node degree matrix

# Model Instance (5)

$$\arg \min_{W_*, b_*} \mathcal{L} = \int_0^T |X(t) - \hat{X}(t)| dt$$

$$\text{subject to } X_h(t) = \tanh(X(t)W_e + b_e)W_0 + b_0$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t)W + b), X_h(0)$$

$$X(t) = X_h(t)W_d + b_d$$

- $W \in \mathbb{R}^{d_e \times d_e}$  and  $b \in \mathbb{R}^{n \times d_e}$ : shared parameters (namely, the weights and bias of a linear connection layer) over time  $t \in [0, T]$
- $W_e \in \mathbb{R}^{d \times d_e}$  and  $W_0 \in \mathbb{R}^{d_2 \times d}$ : for decoding
- $b_e, b_0, b, b_d$ : the biases at the corresponding layer

## Model Instance (6)

$$\arg \min_{W_*, b_*} \mathcal{L} = \int_0^T |X(t) - \hat{X}(t)| dt$$

$$\text{subject to } X_h(t) = \tanh(X(t)W_e + b_e)W_0 + b_0$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t)W + b), X_h(0)$$

$$X(t) = X_h(t)W_d + b_d$$

- We learn the parameters

$$W_e, W_0, W, W_d, b_e, b_0, b, b_d$$

from empirical data so that we can learn  $X$  in a data-driven manner

# Model Instance (7)

$$\arg \min_{W_*, b_*} \mathcal{L} = \int_0^T |X(t) - \hat{X}(t)| dt$$

$$\text{subject to } X_h(t) = \tanh(X(t)W_e + b_e)W_0 + b_0$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t)W + b), X_h(0)$$

$$X(t) = X_h(t)W_d + b_d$$

- $\frac{dX(t)}{dt}$ : a single neural layer at time moment  $t$
- $X(t)$  at arbitrary time  $t$  is achieved by integrating  $\frac{dX(t)}{dt}$  over time, leading to a continuous-time deep neural network:

$$X(t) = X(0) + \int_0^t \text{ReLU}(\Phi X(\tau)W + b) d\tau$$

# Outline

- 1 General Framework
  - Neural Dynamics on Complex Networks (NDCN)
- 2 Learning Continuous-Time Network Dynamics
  - Model Instance
  - Experiments
- 3 Learning Regularly-Sampled Dynamics
  - Baselines, Experimental Setup and Results
- 4 Learning Semantic Labels at Terminal Time
  - Model Instance
  - Experiments



# Baselines

- There are no baselines for learning continuous-time dynamics on complex networks
- Thus we compare the ablation models of NDCN
- By investigating ablation models we show that NDCN is a minimum model for this task

# Baseline #1

- Keep the loss function the same
- The model without encoding and decoding functions
- Thus no hidden space:

$$\frac{dX(t)}{dt} = \text{ReLU}(\Phi X(t)W + b),$$

- Namely ODE-GNN, which learns the dynamics in the original signal space  $X(t)$

## Baseline #2

- Keep the loss function the same
- The model without graph diffusion operator

$$\Phi : \frac{dX_h(t)}{dt} = \text{ReLU}(X_h(t)W + b),$$

- I.e. an ODE Neural Network, which can be thought as a continuous-time version of forward residual neural network

# Baseline #3

- Keep the loss function the same
- The model without control parameters  $W$ :

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t))$$

- No linear connection layer between  $t$  and  $t + dt$  where  $dt \rightarrow 0$
- Thus indicating a determined dynamics to spread signals

# Experimental Setup (1)

- We generate underlying networks with 400 nodes by Network Dynamics #1-#3 and Complex Networks #1-#5
- We set the initial value  $X(0)$  the same for all the experiments
- Thus different dynamics are only due to their different dynamic rules and underlying networks

## Experimental Setup (2)

- We irregularly sample 120 snapshots of the continuous-time dynamics

$$\{X(\hat{t}_1), \dots, X(\hat{t}_{120}) : 0 \leq t_1 < \dots < t_{120} \leq T\}$$

where the time intervals between  $t_1, \dots, t_{120}$  are different

- Training: Randomly choose 80 snapshots from  $X(\hat{t}_1)$  to  $X(\hat{t}_{100})$
- Interpolation testing: the left 20 snapshots from  $X(\hat{t}_1)$  to  $X(\hat{t}_{100})$
- Extrapolation testing: use the 20 snapshots from  $X(\hat{t}_{101})$  to  $X(\hat{t}_{120})$

## Experimental Setup (3)

- We use Dormand-Prince method to get the ground truth dynamics
- We use Euler method in the forward process of our NDCN
- We evaluate the results by  $\ell_1$  loss and normalized  $\ell_1$  loss (normalized by the mean element-wise value of  $\hat{X}(t)$ ) and they lead to the same conclusion
- Results are the mean and standard deviation of the loss over 20 independent runs for 3 dynamic laws on 5 different networks by each method

# Results (Visual)

- We find that one dynamic law may behave quite different on different networks
  - Heat dynamics may gradually die out to be stable but follow different dynamic pattern on different networks
  - Gene dynamics are asymptotically stable on grid but unstable on random networks or community networks
  - Both gene regulation dynamics and biological mutualistic dynamics show very bursty patterns on power-law networks
- NDCN learns all these different network dynamics very well



# Results (Quantitative)

- Each quantitative result is the normalized  $\ell_1$  error with standard deviation (in percentage %) from 20 runs for 3 dynamics on 5 networks by each method
- NDCN captures different dynamics on various complex networks accurately
- NDCN outperforms all the continuous-time baselines by a large margin
- NDCN potentially serves as a minimum model in learning continuous-time dynamics on complex networks

# Outline

- 1 General Framework
  - Neural Dynamics on Complex Networks (NDCN)
- 2 Learning Continuous-Time Network Dynamics
  - Model Instance
  - Experiments
- 3 Learning Regularly-Sampled Dynamics
  - Baselines, Experimental Setup and Results
- 4 Learning Semantic Labels at Terminal Time
  - Model Instance
  - Experiments

# Baselines

- We compare our model with the temporal-GNN models
  - Temporal-GNN are usually combinations of RNN models and GNN models
  - Temporal-GNN models are usually used for next few step prediction and cannot be used for interpolation task (say, to predict  $X[t_{1.23}]$ )
- We use GCN as a graph structure extractor
- We use LSTM/GRU/RNN to learn the temporal relationship between ordered structured sequences

# Baseline #1

- We keep the loss function the same
- LSTM-GNN: the temporal-GNN with LSTM cell

$$X[t + 1] = \text{LSTM}(\text{GCN}(X[t], G))$$

## Baseline #2

- We keep the loss function the same
- GRU-GNN: the temporal-GNN with GRU cell

$$X[t + 1] = \text{GRU}(\text{GCN}(X[t], G))$$

# Baseline #1

- We keep the loss function the same
- RNN-GNN: the temporal-GNN with RNN cell

$$X[t + 1] = \text{RNN}(\text{GCN}(X[t], G))$$

## Experimental Setup

- We regularly sample 100 snapshots of the continuous-time network dynamics

$$\{X[\hat{t}_1], \dots, X[\hat{t}_{100}] : 0 \leq t_1 < \dots < t_{120} \leq T\}$$

where the time intervals between  $t_1, \dots, t_{100}$  are the same

- Training: use first 80 snapshots  $X[\hat{t}_1], \dots, X[\hat{t}_{80}]$
- Prediction/Extrapolation Testing: use the left 20 snapshots  $X[\hat{t}_{81}], \dots, X[\hat{t}_{100}]$
- We use 5 and 10 for hidden dimension of GCN and RNN models respectively

# Results

- GRU-GNN model works well in mutualistic dynamics on random network and community network
- NDCN predicts different dynamics on these complex networks accurately
- NDCN outperforms the baselines in almost all the settings
- NDCN captures the structure and dynamics in a much more succinct way
- NDCN only has 901 parameters to learn, compared to  $24k$ ,  $64k$ ,  $84k$  of RNN-GCN, GRU-GNN, LSTM-GNN, respectively



# Outline

- 1 General Framework
  - Neural Dynamics on Complex Networks (NDCN)
- 2 Learning Continuous-Time Network Dynamics
  - Model Instance
  - Experiments
- 3 Learning Regularly-Sampled Dynamics
  - Baselines, Experimental Setup and Results
- 4 Learning Semantic Labels at Terminal Time
  - Model Instance
  - Experiments

# Learning the Semantic Labels at the Terminal Time

- Existing GNNs (s.o.t.a. in graph semi-supervised classification task) usually adopt 1 or 2 hidden layers
- NDCN follows the perspective of a dynamical system and goes beyond an integer number  $L$  of hidden layers in GNNS to a real number depth  $t$  of hidden layers, implying continuous-time dynamics on the graph
- By integration continuous-time dynamics on the graph over time, we get a more fine-grained forward process
- Thus NDCN model shows very competitive even better results compared with s.o.t.a. GNN models which may have sophisticated parameters (e.g. attention)

# Model Instance (1)

$$\arg \min_{W_e, b_e, W_d, b_d} \mathcal{L} = \int_0^T \mathcal{R}(t) dt - \sum_{i=1}^n \sum_{k=1}^c \hat{Y}_{i,k}(T) \log Y_{i,k}(T)$$

$$\text{subject to } X_h(0) = \tanh(X(0)W_e + b_e)$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t))$$

$$X(T) = \text{Softmax}(X_h(T)W_d + b_d)$$

- Loss: terminal semantic loss  $\mathcal{S}(Y(T))$  modeled by the cross-entropy loss for classification task

## Model Instance (2)

$$\arg \min_{W_e, b_e, W_d, b_d} \mathcal{L} = \int_0^T \mathcal{R}(t) dt - \sum_{i=1}^n \sum_{k=1}^c \hat{Y}_{i,k}(T) \log Y_{i,k}(T)$$

$$\text{subject to } X_h(0) = \tanh(X(0)W_e + b_e)$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t))$$

$$X(T) = \text{Softmax}(X_h(T)W_d + b_d)$$

- $Y(T) \in \mathbb{R}^{n \times c}$ : the label distributions of nodes at time  $T \in \mathbb{R}$  whose
  - $Y_{i,k}(T)$ : the probability of the node  $i = 1, \dots, n$  with label  $k = 1, \dots, c$  at time  $T$
- $\hat{Y}(T) \in \mathbb{R}^{n \times c}$ : the supervised information (again missing information can be padded by 0) observed at  $t = T$

# Model Instance (3)

$$\arg \min_{W_e, b_e, W_d, b_d} \mathcal{L} = \int_0^T \mathcal{R}(t) dt - \sum_{i=1}^n \sum_{k=1}^c \hat{Y}_{i,k}(T) \log Y_{i,k}(T)$$

$$\text{subject to } X_h(0) = \tanh(X(0)W_e + b_e)$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t))$$

$$X(T) = \text{Softmax}(X_h(T)W_d + b_d)$$

- We use differential equation system  $\frac{dX(t)}{dt} = \text{ReLU}(\Phi X(t))$  to spread the graph signals over continuous time  $[0, T]$ , i.e.,

$$X_h(T) = X_h(0) + \int_0^T \text{ReLU}(\Phi X_h(t))$$

# Model Instance (4)

$$\arg \min_{W_e, b_e, W_d, b_d} \mathcal{L} = \int_0^T \mathcal{R}(t) dt - \sum_{i=1}^n \sum_{k=1}^c \hat{Y}_{i,k}(T) \log Y_{i,k}(T)$$

$$\text{subject to } X_h(0) = \tanh(X(0)W_e + b_e)$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t))$$

$$X(T) = \text{Softmax}(X_h(T)W_d + b_d)$$

- Compared with the continuous-time model instance, we only have supervised information from one snapshot at time  $t = T$
- Thus we model the running loss as the  $\ell_2$ -norm regularizer of the learnable parameters to avoid overfitting:

$$\int_0^T \mathcal{R}(t) dt = \lambda(|W_e|_2^2 + |b_e|_2^2 + |W_d|_2^2 + |b_d|_2^2)$$

# Model Instance (5)

$$\arg \min_{W_e, b_e, W_d, b_d} \mathcal{L} = \int_0^T \mathcal{R}(t) dt - \sum_{i=1}^n \sum_{k=1}^c \hat{Y}_{i,k}(T) \log Y_{i,k}(T)$$

$$\text{subject to } X_h(0) = \tanh(X(0)W_e + b_e)$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t))$$

$$X(T) = \text{Softmax}(X_h(T)W_d + b_d)$$

- We adopt the diffusion operator  $\Phi = \tilde{D}^{-\frac{1}{2}}(\alpha I + (1 - \alpha)A)\tilde{D}^{-\frac{1}{2}}$  where  $A$  is the adjacency matrix,  $D$  is the degree matrix and  $\tilde{D} = \alpha I + (1 - \alpha)D$  keeps  $\Phi$  normalized

## Model Instance (6)

$$\arg \min_{W_e, b_e, W_d, b_d} \mathcal{L} = \int_0^T \mathcal{R}(t) dt - \sum_{i=1}^n \sum_{k=1}^c \hat{Y}_{i,k}(T) \log Y_{i,k}(T)$$

$$\text{subject to } X_h(0) = \tanh(X(0)W_e + b_e)$$

$$\frac{dX_h(t)}{dt} = \text{ReLU}(\Phi X_h(t))$$

$$X(T) = \text{Softmax}(X_h(T)W_d + b_d)$$

- The parameter  $\alpha \in [0, 1]$  tunes nodes' adherence to their previous information or their neighbors' collective opinion
- We use  $\alpha$  as a hyper-parameter here for simplicity and we can make it as a learnable parameter later



## Model Instance (7)

- The differential equation system  $\frac{dX}{dt} = \Phi X$  follows the dynamics of averaging the neighborhood opinion as

$$\frac{d\overrightarrow{x_i(t)}}{dt} = \frac{\alpha}{(1-\alpha)d_i + \alpha} \overrightarrow{x_i(t)} + \sum_j^n A_{i,j} \frac{1-\alpha}{\sqrt{(1-\alpha)d_i + \alpha} \sqrt{(1-\alpha)d_j + \alpha}} \overrightarrow{x_j(t)}$$

for node  $i$

- When  $\alpha = 0$ ,  $\Phi$  averages the neighbors as normalized random walk
- When  $\alpha = 1$ ,  $\Phi$  captures exponential dynamics without network effects
- When  $\alpha = 0.5$ ,  $\Phi$  averages both neighbors and itself

# Outline

- 1 General Framework
  - Neural Dynamics on Complex Networks (NDCN)
- 2 Learning Continuous-Time Network Dynamics
  - Model Instance
  - Experiments
- 3 Learning Regularly-Sampled Dynamics
  - Baselines, Experimental Setup and Results
- 4 Learning Semantic Labels at Terminal Time
  - Model Instance
  - Experiments

# Results (1)

- NDCN outperforms many s.o.t.a. GNN models
- We report the mean and standard deviation of our results for 100 runs
- Cora dataset: terminal time  $T = 1.2, \alpha = 0$
- Citeseer dataset:  $T = 1.0, \alpha = 0.8$
- Pubmed dataset:  $T = 1.1, \alpha = 0.4$

## Results (2)

- NDCN gives better classification accuracy at terminal time  $T \in \mathbb{R}^+$  by capturing the continuous-time network dynamics to diffuse network signals
- For all the three datasets their accuracy curves follow rise and fall patterns around the best terminal time
- When the terminal time  $T$  is too small or too large, the accuracy degenerates because the features of nodes are in under-diffusion or over-diffusion states, implying the necessity in capturing continuous-time dynamics
- In contrast, previous GNNs can only have a discrete number of layers which cannot capture the continuous-time network dynamics accurately